

Can you use a hashtable to implement skipTo()?


Better than next()

- What's the worst case for sequential merge-based intersection?
- {52, 1} → **Assignment Project Exam Help**
<https://powcoder.com>
– To the position whose id is at least 52 → **skipTo(52)**
– Essentially, asking the first i , such that $K_2[i] \geq 52$ (K_2 's list is sorted).
Add WeChat powcoder
– Takes many sequential call of next()
– Could use binary search in the rest of the list
– Cost: $\lceil \log_2(N_{\text{remainder}}) \rceil$

K_2 :	1	3	5	79
K_1 :	52	54	56	58				

skipTo(id)

- Galloping search (gambler's strategy)
 - [Stage 1] Doubling the search range until you overshoot
 - [Stage 2] Perform binary search in the last range
- Performance analysis (worst case)
 - Let the destination position be n positions away.
 - $\approx \log_2 n$ probes in Stage 1 + $\approx \log_2 n$ probes in Stage 2
 - Total = $2 \lceil \log_2 (n+1) \rceil = O(\log_2 n)$



Total Cost

- Galloping search (gambler's strategy)
 - Cost of the i -th probe: $\approx 2 \log_2(n_i)$
 - Total cost of K_1 probes: $\approx 2 \log_2(\prod_{i=1}^{|K_1|} n_i)$
 $\leq 2 \log_2(((\sum_{i=1}^{|K_2|} n_i) / |K_1|)^{|K_1|}) \leq 2|K_1| \log_2(|K_2|/|K_1|)$
- Asymptotically, resembles linear merge when $|K_2|/|K_1| = O(1)$, resembles binary search when $|K_1| = O(1)$

What about list intersection using binary search?

Multiple Term Conjunctive Queries

- K_1 AND K_2 AND ... AND K_n
- SvS does not perform well if none of the associated lists are short
- In addition, it is blocking
- Can you design non-blocking multiple sorted array intersection algorithm?

Generalization

- Generalize the 2-way intersection algorithm

Assignment Project Exam Help

- 2-way: <https://powcoder.com>

– {1, 2} → move k_1 's cursor

– skipTo(2)

K_1 :

1	3
---	---

K_2 :

2	4	6
---	---	---

K_3 :

3	9	27	81
---	---	----	----

- 3-way:

– {1, 2, 3} → move k_1, k_2 's cursor

– skipTo(3)

eliminator = $\text{Max}_{1 \leq i \leq n}(k_i.\text{cursor})$

Optimization

- Mismatch found even before accessing K_3 's cursor

Assignment Project Exam Help

- Choice 1: continue to get cursors of other list

<https://powcoder.com>

Add WeChat powcoder

- Choice 2: settle the

K_1 :	1	3		
K_2 :	2	4	6	
K_3 :	3	9	27	81

dispute within the first two lists → max
algorithm [Culpepper & Moffat, 2010]

- Better locality of access → fewer cache misses
- Similar to SvS

Pseudo-Code for the **Max** Algorithm (Wrong)

- Input: K_1, K_2, \dots, K_n in increasing size

```
(1)  $x := K_1[1]$ ;  $startAt := 2$  //x is the eliminator
(2) while  $x$  is defined do
(3)   for  $i = startAt$  to  $n$  do
(4)      $y := K_i.skipTo(x)$ 
(5)     if  $y > x$  then //mismatch
(6)        $x := K_1.next()$  //restart_1 //restart_2
(7)       if  $y > x$  then  $startAt := 1$   $x := y$  else  $startAt := 2$  end if
(8)       break //match in all lists
(9)     elsif  $i = n$  then //y = x
(10)      Output  $x$ 
(11)       $x := K_1.next()$ 
(12)    end if
(13)  end for
(14) end while
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

A



B



Assignment Project Exam Help

<https://powcoder.com>

C



Add WeChat powcoder

D



The original code has a bug when in restart_1 cases

Pseudo-Code for the **Max** Algorithm (Fixed)

- Input: K_1, K_2, \dots, K_n in increasing size

```
(1)  $x := K_1[1]$ ;  $startAt := 2$ 
(2) while  $x$  is defined do
(3)     for  $i = startAt$  to  $n$  do
(4)          $y := K_i.skipTo(x)$ 
(5)         if  $y > x$  then
(6)              $x := K_1.next()$ 
(7)             if  $y > x$  then  $startAt := 1$  else  $startAt := 2$  end if
(8)             break
(9)         elseif  $i = n$  then
(10)            Output  $x$ 
(11)             $x := K_1.next()$ 
(12)        end if
(13)    end for
(14) end while
```

(4.1) **if** $i = 1$ **then**
(4.2) **if** $y > x$ **then**
(4.3) $x := y$
(4.4) **break**
(4.5) **end if**
(4.6) **end if**

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

References

- J. Shane Culpepper, Alistair Moffat: Efficient set intersection for projected indexing. ACM Trans. Inf. Syst. 29(1): 1 (2010)
<https://powcoder.com>
- F.K. Hwang and S. Lin, A simple algorithm for merging two disjoint linearly ordered sets. SIAM J. Comput. 1 1 (1972), pp. 31–39.
Add WeChat powcoder
- Stefan Buettcher, Charles L. A. Clarke, Gordon V. Cormack, Information Retrieval: Implementing and Evaluating Search Engines, 2010 [Chapter 5]